

# Simulation analysis

leios-2025w23

leios-2025w24

# Experiments

1. Pseudo-mainnet (`leios-2025w23`)
  - a. Transaction lifecycle (Rust only)
2. Mini-mainnet (`leios-2025w24`)
  - a. Transaction lifecycle (Rust only)
  - b. IB diffusion (comparison of Rust vs Haskell)

## Goals

- Compare network topologies
- Re-check Haskell vs Rust
- Realistic scenarios on mini-mainnet

# Pseudo-mainnet

This is the first cut at a realistic mainnet-scale topology for Leios, but it likely contain imperfections because several compromises were made during its construction, so as to smooth out inconsistencies in source data.

- Realistic stake distribution
- Realistic number of stake pools
- Two relays for each block producer
- Block producers only connected to their relays
- 10,000 nodes total
- Realistic latencies, generally consistent with the RIPE Atlas ping dataset
- Bandwidth consistent with the low end of what is generally available in cloud data centers
- Node connectivity generally consistent with measurements by the Cardano Foundation
- Geographic distribution (countries and autonomous systems) consistent with measurements by the Cardano Foundation

**Finding:** *Creating a much more realistic topology would require simulating the p2p algorithm itself at mainnet scale, so that the topology would be an emergent property of the simulation.*

Metric	Value
Total nodes	10000
Block producers	2657
Relay nodes	7343
Total connections	298756
Network diameter	6 hops
Average connections per node	29.88
Clustering coefficient	0.122
Average latency	77.0 ms
Maximum latency	636.8 ms
Stake-weighted latency	0.0 ms
Bidirectional connections	10800
Asymmetry ratio	92.77%

# Mini-mainnet

- The “mini-mainnet” was created using the same procedure and data sources as for pseudo-mainnet.
- However, . . .
  - the smaller stakepools were eliminated, and
  - the total number of nodes was reduced to 750.
- The network has a slightly smaller diameter, but its other characteristics are similar to pseudo-mainnet.
- Whereas running simulations on pseudo-mainnet was nearly impractical, running them on mini-mainnet is eminently feasible.

Metric	Value
Total nodes	750
Block producers	216
Relay nodes	534
Total connections	19314
Network diameter	5 hops
Average connections per node	25.75
Clustering coefficient	0.332
Average latency	64.8ms ms
Maximum latency	578.3ms ms
Stake-weighted latency	0.0ms ms
Bidirectional connections	1463
Asymmetry ratio	84.85%


# Transaction lifecycle experiments

- Rust simulator (since Haskell does not model transactions)
- 100-node topology
- 8 vCPUs / node
- 10 slot / stage
- 3 shard / group
- 10 groups
- 1.5 EB/stage
- 1, 3, 10, 30, 100, 300 tx/s\*
- IB generation probability varies with TPS
- Number of shards  $\approx (\text{IB rate}) * (\text{shards per group}) * (\text{groups})$
- 327,680 B/IB maximum
- **No unsharded transactions**


Comparing pseudo-mainnet to mini-mainnet

# Comparison of message diffusion

## *Pseudo-mainnet*



## *Mini-mainnet*



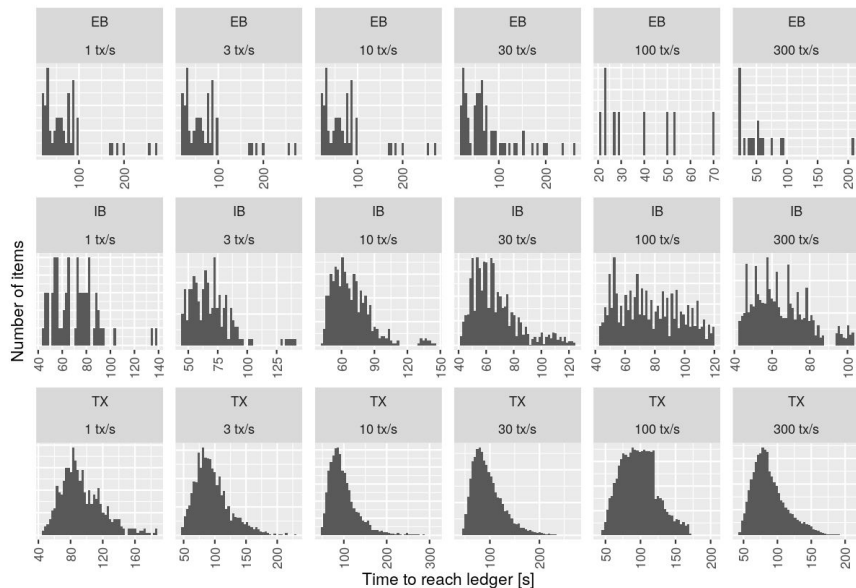
Arrival-time distributions are similar, but pseudo-mainnet has more dispersion.

# Comparison of time to reach the ledger

## *Mini-mainnet*

Time to reach the ledger

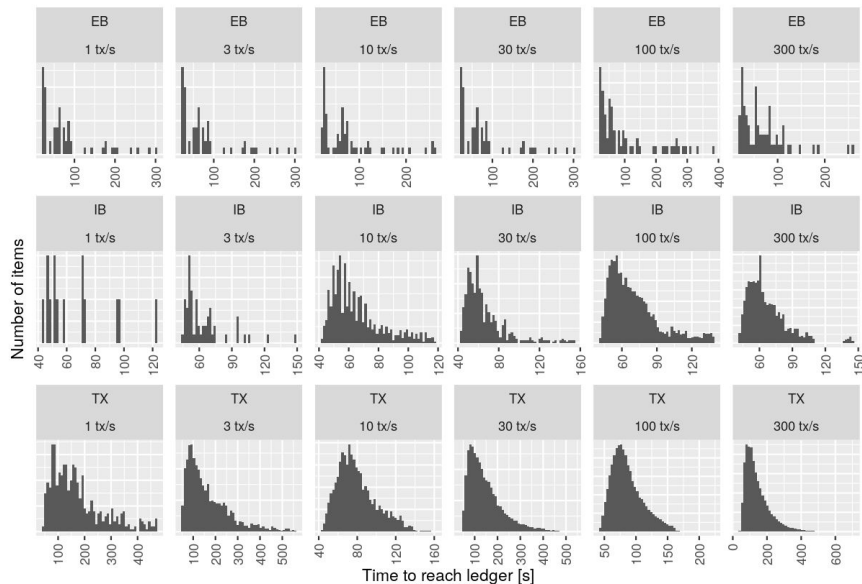
Rust simulator, pseudo-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards



## *Pseudo-mainnet*

Time to reach the ledger

Rust simulator, mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards




Ledger-time distributions are similar, but pseudo-mainnet is slightly faster.

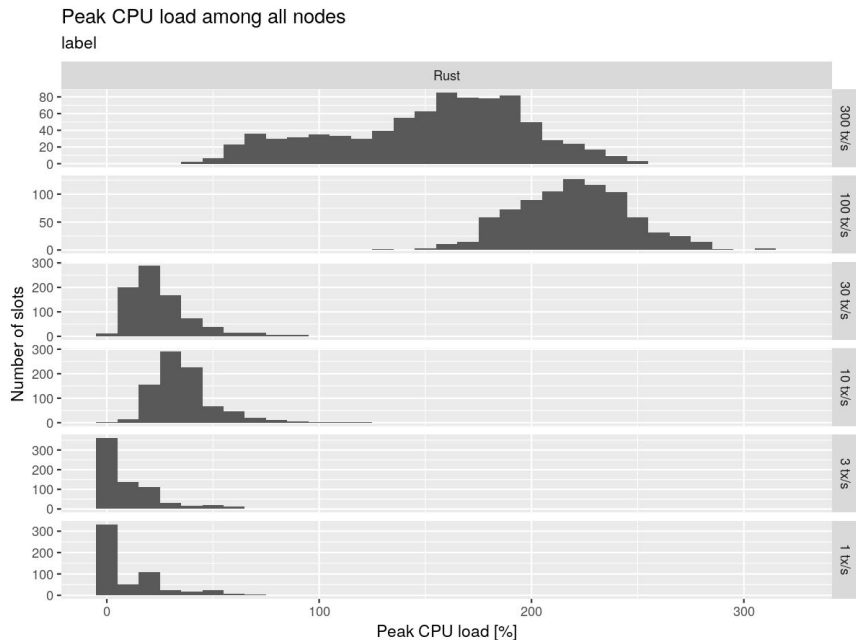


# Comparison of peak CPU usage

## *Mini-mainnet*



## *Pseudo-mainnet*



Peak CPU usage is quite similar, except in the 100 TPS case.

# Findings about mainnet variants


- Qualitative comparison of key metrics indicates little substantial difference between simulation results using pseudo-mainnet vs mini-mainnet.
  - Mini-mainnet might actually be more stressful to Leios than pseudo-mainnet.
- Differences would not affect recommendations about Leios.
  - Similar resource requirements
  - Similar performance metrics
- Future Leios simulations will use mini-mainnet because it is far more computationally tractable.
  - Pseudo-mainnet simulations might be used again at the close of the innovation stream, in order to validate any conclusions based on mini-mainnet simulations.
- Bandwidth assumptions are more important than topology assumptions.
- Next-generation Leios simulators should implement dynamic p2p faithfully, so as to eliminate the expedient of simulation with a static topology.

# Comparing Rust and Haskell simulators on mini-mainnet

# Comparison of message diffusion


Arrival delay for IB

mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards



Arrival delay for EB


mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards



Diffusion of messages is consistent between Rust and Haskell, except for IBs at 32 IB/s, where Haskell is more diffuse.


Arrival delay for VT

mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards



Arrival delay for RB


mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards



# Comparison of IB delivery

Arrival fraction for IB

mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards



Both Haskell and Rust struggle above 15 IB/s to deliver all IBs, but Haskell struggles more.


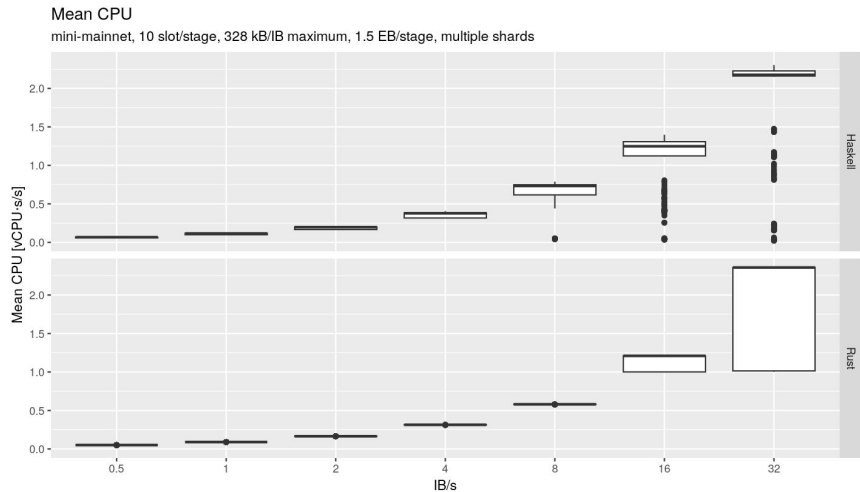
This IB rate corresponds to

- ~50 MB/s
- ~900X Praos
- > 10k TPS

Caveats:

- 1 Gb/s node-to-node links
- IBs are full but TXs are not diffused

# Comparison of CPU usage



Rust shows a broader spread in CPU usage among nodes, but both simulations have similar maxima.

# Comparison of CPU tasks

Mean CPU load among all nodes

mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards



Overall CPU load is comparable between Haskell and Rust, but Haskell appears to spend more time on IB header validation.

# Findings about simulator variants

- Haskell and Rust behave similarly on the mini-mainnet topology.
  - Differences are moderately significant only at high rates like 32 IB/s.
- Future simulations studies will rely on the Rust simulator because it is much faster and models transactions.
  - The Haskell simulator might be used again at the close of the innovation stream, in order to re-validate any conclusions based on the Rust simulator.



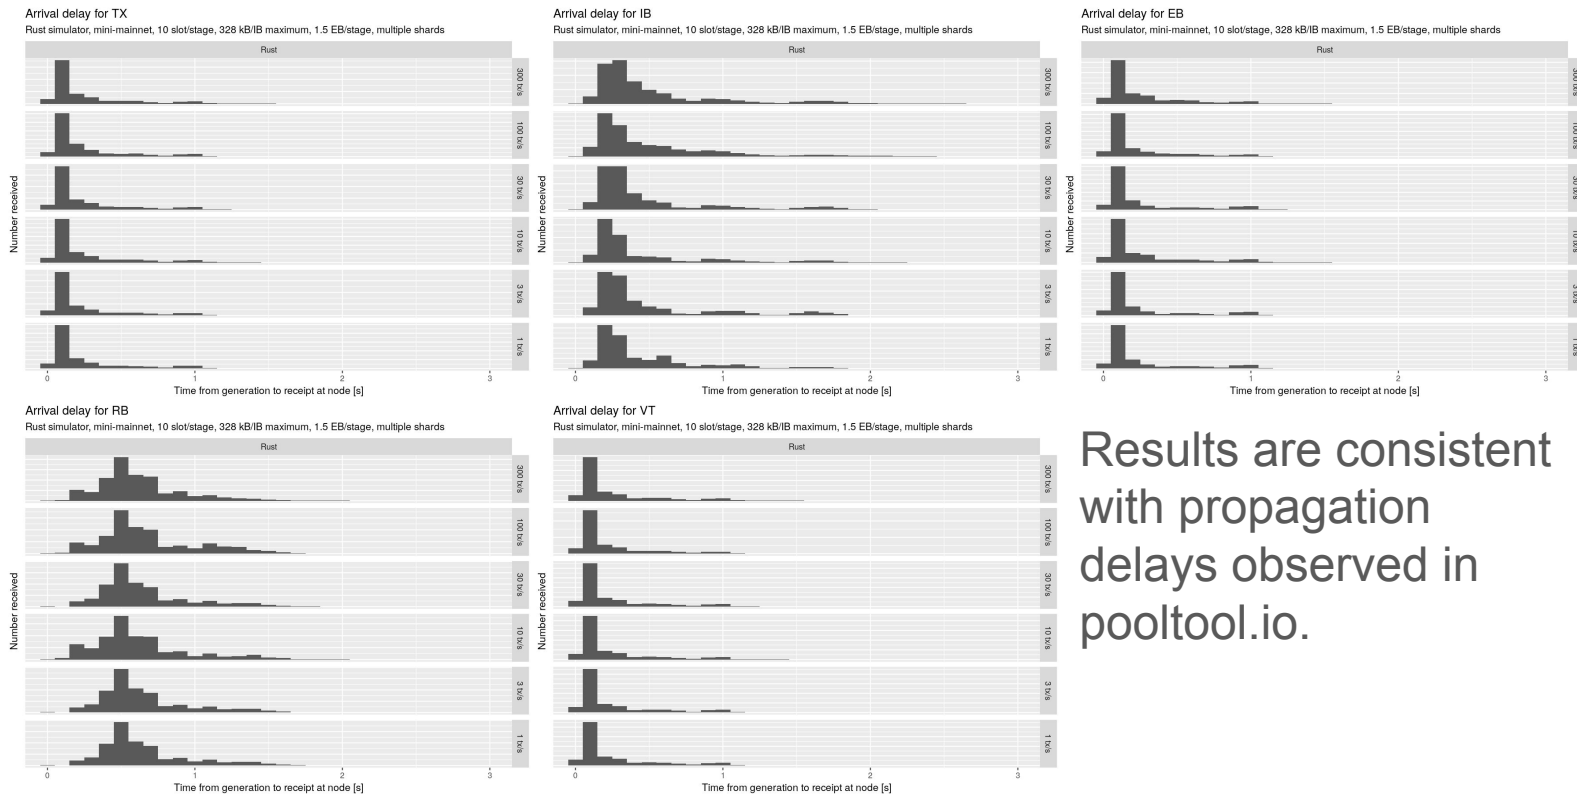
Rust results on mini-mainet at  
leios-2025w24

# Performance metrics

*Performance is stable to 300+ TPS.*

Nominal throughput	TX in IB	TX per IB	IB per EB	IB in EB	Spatial efficiency	TX redundancy	Time to IB	Time to EB	Time to ledger
1 TPS	1.071	31.579	0.851	5.421	67.187%	79.717%	107.129s	131.667s	173.349s
3 TPS	1.034	22.612	3.777	5.506	85.475%	74.333%	94.802s	124.359s	166.205s
10 TPS	1.203	3.002	120.292	5.945	73.887%	81.706%	13.040s	13.040s	79.738s
30 TPS	1.010	22.843	37.397	5.839	94.259%	77.360%	80.816s	109.165s	147.285s
100 TPS	1.262	3.371	1179.405	6.127	72.659%	82.428%	13.336s	41.658s	86.530s
300 TPS	1.027	27.059	386.429	5.692	91.073%	79.845%	86.223s	111.467s	144.033s


# Message arrival (with implications for concurrency)



## Time for transaction to reach the ledger

Rust simulator, mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards

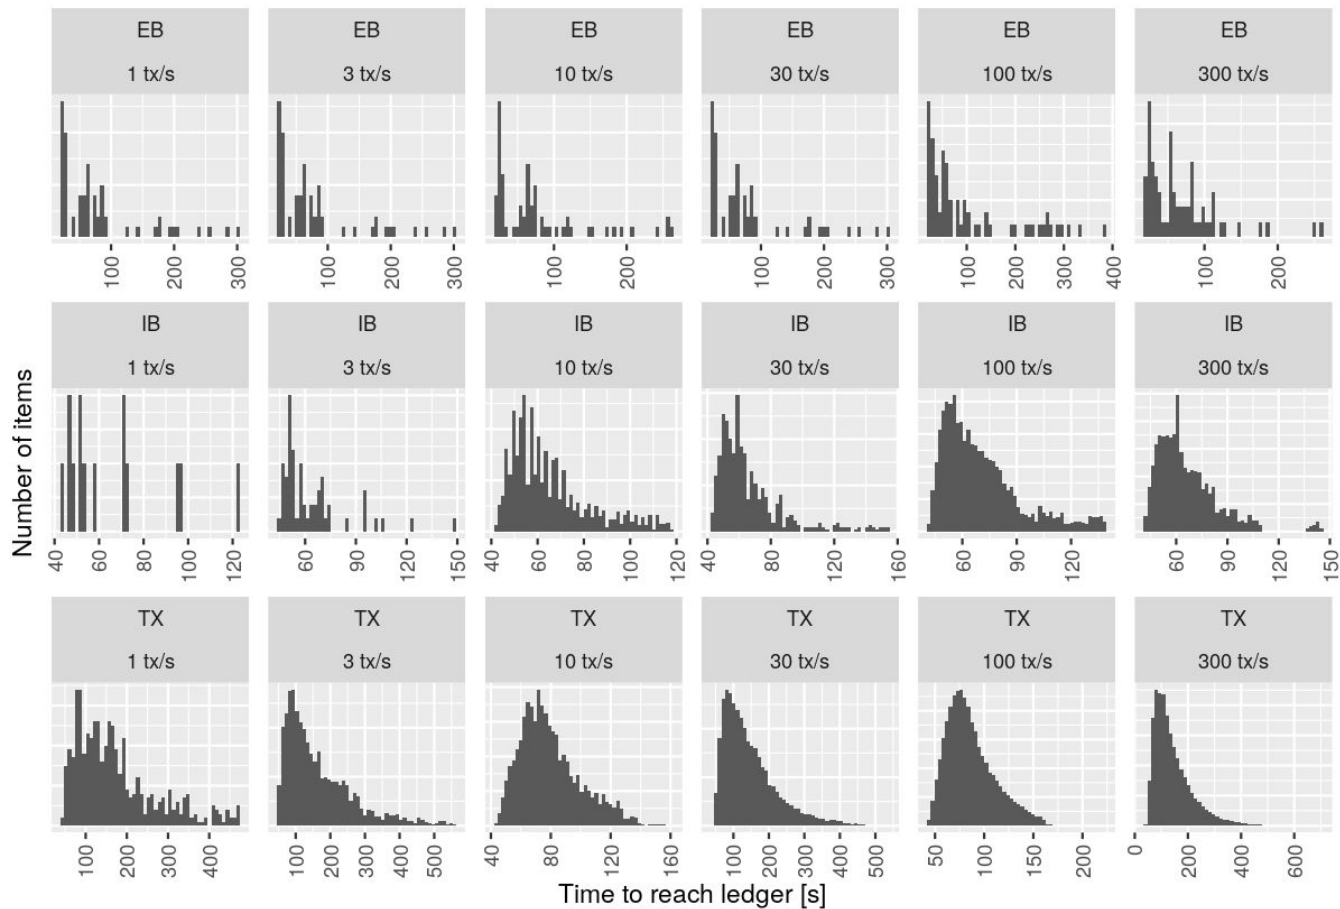
Results are similar to previous simulations, but there is an odd non-monotonicity in the spread of ledger times.



Results are similar to the 100-node topology.

## Time to reach the ledger


Rust simulator, mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards



# Transactions reach the ledger, except when the simulation was stopped too soon.


## Transactions reaching the ledger

Rust simulator, mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards



## Transactions reaching the ledger


Rust simulator, mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards



The total size of IBs is on the order of the size of the transactions.

#### Size of persisted data


Rust simulator, mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards



The amount of persistent storage is somewhat larger than the total size of transactions being stored.

#### Spatial efficiency (size of txs on ledger / size of non-tx persisted data)

Rust simulator, mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards



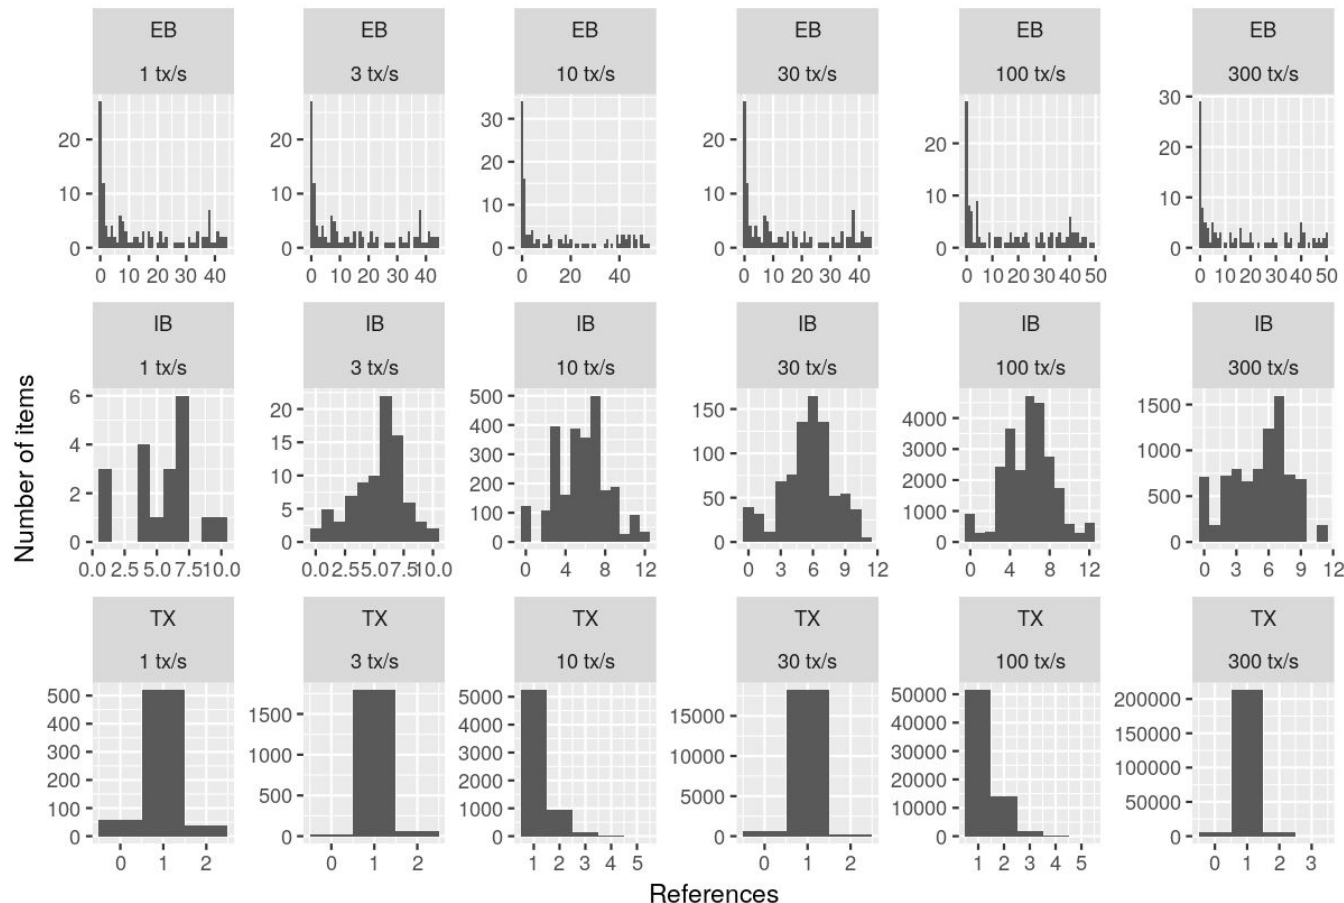
Number of references (0 = not used, 2+ = duplicated)

Rust simulator, mini-mainnet, 10 slot/stage, 328 kB/IB maximum, 1.5 EB/stage, multiple shards

EBs are often referenced by other EBs multiple times.

IBs are often referenced by EBs multiple times.

Transactions are occasionally included in multiple IBs.





# Findings from leios-2025w24 simulations

- The refactored simulation workflow allows rapid processing of large simulation results.
- The mini-mainnet topology balances realism with speed of simulation.
- We need better empirical evidence for the effective bandwidth of node-to-node connections because this sets the fundamental limit on Leios performance.
- Once the candidate Leios variants are finalized, we can run additional experiments:
  - Comparison of variants
  - Optimal protocol parameters
  - Performance measurements for inclusion in the CIP